

Hierarchical Memory & Cache Simulator

Architectural Design Document

System Implementation Report

Contents

1	Introduction	1
2	System Architecture & Assumptions	1
2.1	Memory Layout	1
2.2	Operational Assumptions	1
3	Component Implementation	1
3.1	Physical Memory Allocator	1
3.1.1	Free List Allocator	1
3.1.2	Buddy System Allocator	2
3.2	Virtual Memory Unit (MMU)	2
3.3	Cache Hierarchy	2
4	Integration & Data Flow	2
5	Limitations & Constraints	3

1 Introduction

This document delineates the architectural specifications and design rationale for a systems-level memory simulator. The project models the critical data path of a computer architecture, encompassing the Central Processing Unit (CPU), Memory Management Unit (MMU), a three-level Cache Hierarchy, and Physical Main Memory.

The primary objective is to facilitate runtime analysis of memory management strategies, providing granular metrics on internal/external fragmentation, page fault latency, and cache hit/miss propagation.

2 System Architecture & Assumptions

2.1 Memory Layout

The simulation operates within a deterministic environment defined by the following constraints:

- **Address Space:** A 32-bit addressing scheme is employed, utilizing `uint32_t` for all virtual and physical addresses.
- **Physical Memory (RAM):** Modeled as a contiguous byte array. The default configuration allocates 4096 Bytes, modifiable at runtime.
- **Page Size:** Fixed at 64 Bytes, establishing the granularity for frames and page table entries.
- **Granularity:** Memory access occurs at the byte level, though the Buddy System allocator enforces power-of-2 alignment.

2.2 Operational Assumptions

To maintain simulation feasibility, the design incorporates the following abstractions:

- **Single Execution Context:** The system models a distinct, single-threaded process. Context switching and thread concurrency are outside the current scope.
- **Symbolic I/O:** Disk storage is abstracted. Page faults trigger a symbolic latency penalty (e.g., 10,000 cycles) rather than performing physical I/O operations.
- **Unified Cache Model:** The cache hierarchy treats instruction fetches and data operands identically, utilizing a unified instruction/data cache structure.

3 Component Implementation

3.1 Physical Memory Allocator

Memory management is handled through two interchangeable allocation paradigms, utilizing the Strategy Design Pattern to permit runtime switching.

3.1.1 Free List Allocator

This implementation utilizes a Doubly Linked List to track free memory segments. Each block contains a metadata header defining its size and allocation status. Supported placement strategies include:

- **First Fit:** Traverses the list linearly and selects the first block sufficient for the request.

- **Best Fit:** Exhaustively searches for the block yielding the smallest residual hole, minimizing internal fragmentation.
- **Worst Fit:** Selects the largest available block, intended to maximize the size of remaining free chunks.

3.1.2 Buddy System Allocator

This allocator manages memory via an implicit binary tree of power-of-2 sized blocks.

- **Splitting:** Allocation requests are rounded up to the nearest power of 2. If a block is too large, it is recursively bisected until the optimal size is reached.
- **Coalescing:** Upon deallocation, the system identifies the block's "buddy" (calculated via bitwise XOR). If the buddy is free, the two blocks merge recursively into a higher-order block.

3.2 Virtual Memory Unit (MMU)

The MMU decouples the virtual address space from physical memory.

- **Page Table:** Implemented via a hashed map structure ($\text{VPN} \rightarrow \text{PageEntry}$) to simulate sparse address spaces efficiently.
- **Replacement Policies:**
 - **FIFO:** Utilizes a queue to evict the oldest loaded page.
 - **LRU:** Maintains a list where successful accesses promote pages to the Most Recently Used (MRU) position, ensuring commonly accessed data remains resident.

3.3 Cache Hierarchy

The system employs a 3-level inclusive cache hierarchy to simulate memory latency.

Level	Size / Assoc	Latency	Role
L1 Cache	128 B / Direct Mapped	1 Cycle	High-speed lookup
L2 Cache	256 B / 2-Way Set	10 Cycles	Intermediate buffer
L3 Cache	512 B / 4-Way Set	100 Cycles	Last Level Cache (LLC)

Table 1: Cache Hierarchy Configuration

Addresses are decomposed into **Tag**, **Index**, and **Offset** fields. A miss at Level N propagates a request to Level $N + 1$, accumulating latency penalties.

4 Integration & Data Flow

The architecture strictly enforces a $\text{Virtual} \rightarrow \text{Physical} \rightarrow \text{Cache}$ translation pipeline. The following diagram illustrates the lifecycle of a CPU memory request.

Execution Sequence:

1. **Translation:** The Virtual Address is split into a Virtual Page Number (VPN) and Offset. The Page Table translates the VPN to a Physical Frame Number.
2. **Fault Handling:** If translation fails, a Page Fault is raised. The policy engine evicts a victim page (if necessary), loads the required data, and updates the Page Table.




Figure 1: Memory Access Data Path

3. **Cache Access:** The resulting Physical Address is passed to the cache controller.

4. **Latency Accumulation:** Total cycle count is aggregated:

$$T_{total} = T_{translation} + T_{L1} + (Miss?T_{L2} : 0) + \dots + (Miss?T_{RAM} : 0).$$

5 Limitations & Constraints

- **Data Persistence:** The cache hierarchy stores metadata (Tags/Valid Bits) to simulate hit/miss behavior but does not replicate the actual data payload.
- **Direct Translation:** A Translation Lookaside Buffer (TLB) is not modeled; all translations incur the full Page Table lookup cost.
- **Thread Safety:** The simulator is designed for sequential execution and is not thread-safe.